

Python Feature Engineering Cookbook

Stephen CUI¹

2022 年 12 月 18 日

¹cuixuanStephen@gmail.com

目录

1	Imputing Missing Data	3
1.1	Technical requirements	3
1.2	Removing observations with missing data	4
1.2.1	How to do it...	4
1.3	Performing mean or median imputation	6
1.3.1	How to do it...	6
1.3.2	How it works...	8
1.4	Imputing categorical variables	8
1.4.1	How to do it...	8
1.4.2	How it works...	10
1.4.3	Replacing missing values with an arbitrary number	11
1.4.4	How to do it...	11
1.5	Finding extreme values for imputation	12
1.5.1	How to do it...	13
1.5.2	How it works...	14
1.6	Marking imputed values	14
1.6.1	How to do it...	15
1.6.2		17

Chapter 1

Imputing Missing Data

Missing data, that is, the absence of values for certain observations, is an unavoidable problem in most data sources.

The act of replacing missing data with statistical estimates of missing values is called imputation. The goal of any imputation technique is to produce a complete dataset. There are multiple imputation methods that we can use, depending on whether the data is missing at random, the proportion of missing values, and the machine learning model we intend to use.

1.1 Technical requirements

We will use the Credit Approval dataset from the UCI Machine Learning Repository (<https://archive.ics.uci.edu/>). To prepare the dataset, follow these steps:

1. Visit <https://archive.ics.uci.edu/ml/machine-learning-databases/credit-screening/>.
2. Click on **crx.data** to download the data.

```
1 import random
2 import numpy as np
3 import pandas as pd
4 data = pd.read_csv('data/crx.data', header=None)
5 varname = [f"A{s}" for s in range(1, 17)]
6 data.columns = varname
7 data = data.replace("?", np.nan)
8 data["A2"] = data["A2"].astype("float")
9 data["A14"] = data["A14"].astype("float")
10 data["A16"] = data["A16"].map({"+": 1, "-": 0})
11 data.rename(columns={"A16": "target"}, inplace=True)
```

```
12 random.seed(37)
13 values = list(set([random.randint(0, len(data)) for p in range(0, 100)]))
14 data.loc[values, ["A3", "A8", "A9", "A10"]] = np.nan
15 data.to_csv('data/credit_approval_uci.csv', index=False)
```

Tip

Make sure you store the dataset in the same folder from which you will execute the code in the recipes.

处理后的数据集有一个target字段以及15个（A1-A15）特征字段组成。其中A11, A12, A13, A15, Target不包含缺失值，Figure 1.1。

1.2 Removing observations with missing data

Complete Case Analysis (CCA), also called list-wise deletion of cases, consists of discarding observations with missing data. CCA can be applied to both categorical and numerical variables. With CCA, we preserve the distribution of the variables after the imputation, provided the data is missing at random and only in a small proportion of observations. However, if data is missing for many variables, CCA may lead to the removal of a large portion of the dataset.

1.2.1 How to do it...

```
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 data = pd.read_csv('data/credit_approval_uci.csv')
4 with plt.style.context('ggplot'):
5     data.isnull().mean().sort_values(ascending=True).plot.bar(rot=45)
6     plt.ylabel("Proportion of missing data")
7     plt.title("Proportion of missing data per variable")
```




图 1.1: Proportion of missing data per variable

```
1 data_cca = data.dropna()
2 print(f"Total number of observations: {len(data)}")
3 print(f"Number of observations without missing data: {len(data_cca)}")
4 from feature_engine.imputation import DropMissingData
5 cca = DropMissingData(variables=None, missing_only=True)
6 cca.fit(data)
7 print(cca.variables_)
8 data_cca = cca.transform(data)
```

Tip

The `dropna()` method drops observations with any missing value by default. We can remove observations with missing data in a subset of variables like this: `data.dropna(subset=["A3", "A4"])`.

Tip

To remove observations with missing data in a subset of variables, use `DropMissingData(variables=['A3', 'A4'])`. To remove observations with missing values in at least 5% of the variables, use `DropMissingData(threshold=0.95)`.

1.3 Performing mean or median imputation

Mean or median imputation consists of replacing missing values with the mean or median variable. The mean or median is calculated using a train set, and these values are used to impute missing data in train and test sets, as well as in all future data we intend to use with the machine learning model. Scikit-learn and feature-engine transformers learn the mean or median from the train set and store these parameters for future use out of the box.

Tip

Use mean imputation if variables are normally distributed and median imputation otherwise. Mean and median imputation may distort the distribution of the original variables if there is a high percentage of missing data.

1.3.1 How to do it...

```

1  import pandas as pd
2  from sklearn.model_selection import train_test_split
3  from sklearn.impute import SimpleImputer
4  from sklearn.compose import ColumnTransformer
5  from feature_engine.imputation import MeanMedianImputer
6
7  data = pd.read_csv('data/credit_approval_uci.csv')
8
9  X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
10                                                         axis='columns'),
11                                                         data['target'],
12                                                         test_size=.3,
13                                                         random_state=39)
14  numeric_vars = X_train.select_dtypes(exclude='O').columns.to_list()
15  numeric_vars
16  median_values = X_train[numeric_vars].median().to_dict()
17  median_values
18  for df in [X_train, X_test]:
19      # unable to modify X_train and X_test
20      # because df pointer to different address
21      #     df = df.fillna(value=median_values)
22      df.fillna(value=median_values, inplace=True)
23      # equivalent to following codes
24      # X_train = X_train.fillna(value=median_values)
25      # X_test = X_test.fillna(value=median_values)

```

```

26
27 X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
28                                                         axis='columns'),
29                                                         data['target'],
30                                                         test_size=.3,
31                                                         random_state=39)
32 remaining_vars = [var for var in X_test.columns if var not in numeric_vars]
33 remaining_vars
34 imputer = SimpleImputer(strategy='median')
35 ct = ColumnTransformer([('imputer', imputer, numeric_vars)],
36                         remainder='passthrough')
37 ct.fit(X_train)
38 ct.named_transformers_.imputer.statistics_
39 # warning: unable to perform median imputation
40 # for df in [X_train, X_test]:
41 #     df = ct.transform(df)
42 # return NumPy arrays
43 X_train = ct.transform(X_train)
44 X_test = ct.transform(X_test)
45 X_train = pd.DataFrame(X_train, columns=numeric_vars + remaining_vars)
46 X_train
47 X_test = pd.DataFrame(X_test, columns=numeric_vars + remaining_vars)
48 X_test
49
50 X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
51                                                         axis='columns'),
52                                                         data['target'],
53                                                         test_size=.3,
54                                                         random_state=39)
55 imputer = MeanMedianImputer(imputation_method='median', variables=numeric_vars)
56 imputer.fit(X_train)
57 imputer.imputer_dict_
58 X_train = imputer.transform(X_train)
59 X_test = imputer.transform(X_test)
60 print(X_train[numeric_vars].isnull().any().any())
61 print(X_test[numeric_vars].isnull().any().any())

```

The mean or median values should be learned from the train set variables. Thus, we divided the dataset into train and test sets using scikit-learn's `train_test_split()` function. The function takes the predictor variables, the target, the fraction of observations to retain in the test set, and a `random_state` value for reproducibility as arguments. We obtained a train set with 70% of the original observations and a test set with 30% of the original observations. The 70:30 split was done at random.

Tip

`SimpleImputer()` operates on the entire `DataFrame` and returns NumPy arrays. To impute just a subset of variables, we need to use `ColumnTransformer()`. In contrast, `MeanMedianImputer()` can take an entire `DataFrame` and yet it will only impute the specified variables, returning a pandas `DataFrame`.

Categorical variables usually contain strings as values, instead of numbers. We replace missing data in categorical variables with the most frequent category, or with a different string. Frequent categories are estimated using the train set and then used to impute values in the train, test, and future datasets. Thus, we need to learn and store these values, which we can do using `scikit-learn` and `feature-engine`'s out-of-the-box transformers.

[illegible]


```

13                                     random_state=37)
14 # using pandas
15 categorical_vars = X_train.select_dtypes(include='O').columns.to_list()
16 # important: iloc[0]
17 frequent_values = X_train[categorical_vars].mode().iloc[0].to_dict()
18 print(frequent_values)
19
20 X_train = X_train.fillna(value=frequent_values)
21 X_test = X_test.fillna(value=frequent_values)
22 print(X_train[categorical_vars].isnull().any().any())
23 print(X_test[categorical_vars].isnull().any().any())
24
25 # Imputation with a string
26 X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
27                                                         axis='columns'),
28                                                         data['target'],
29                                                         test_size=.3,
30                                                         random_state=37)
31 imputation_dict = {var: 'no_data' for var in categorical_vars}
32 X_train = X_train.fillna(value=imputation_dict)
33 X_test = X_test.fillna(value=imputation_dict)
34 print(X_train['A1'].value_counts())
35
36 # using scikit-learn
37 X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
38                                                         axis='columns'),
39                                                         data['target'],
40                                                         test_size=.3,
41                                                         random_state=37)
42
43 remaining_vars = [
44     var for var in X_train.columns if var not in categorical_vars
45 ]
46
47 imputation_dict = {var: 'not_data' for var in categorical_vars}
48 imputer = SimpleImputer(strategy='most_frequent')
49
50 ct = ColumnTransformer([('Imputer', imputer, categorical_vars)],

```

```

51         remainder='passthrough')
52
53 ct.fit(X_train)
54 # this string "Imputer" here must be identical to
55 # the string 'Imputer' in ColumnTransformer
56 print(ct.named_transformers_.Imputer.statistics_)
57
58 X_train = ct.transform(X_train)
59 X_test = ct.transform(X_test)
60
61 X_train = pd.DataFrame(X_train, columns=categorical_vars + remaining_vars)
62 print(X_train[categorical_vars].isnull().any().any())
63
64 # using feature-engine
65 X_train, X_test, y_train, y_test = train_test_split(data.drop('target',
66                                                         axis='columns'),
67                                                         data['target'],
68                                                         test_size=.3,
69                                                         random_state=37)
70
71 imputer = CategoricalImputer(imputation_method='frequent', variables=categorical_vars)
72
73 imputer.fit(X_train)
74
75 print(imputer.imputer_dict_)
76
77 X_train = imputer.transform(X_train)
78 X_test = imputer.transform(X_test)
79 print(X_train[categorical_vars].isnull().any().any())
80 print(X_test[categorical_vars].isnull().any().any())

```

1.4.2 How it works...

`SimpleImputer()` returns a NumPy array by default. We converted this array into a pandas DataFrame. We had to pass the variable names in the correct order: the imputed variables are located first in the array, followed by the remaining variables. 这个需要特别的注意。

Tip

Note that, unlike `SimpleImputer()`, `CategoricalImputer()` will only impute categorical variables, unless specifically told not to do so by setting the `ignore_format` parameter to `True`.

1.4.3 Replacing missing values with an arbitrary number

When replacing missing values with arbitrary numbers, we need to be careful not to select a value close to the mean, the median, or any other common value of the distribution.

Tip

Arbitrary number imputation can be used when data is not missing at random, when we are building non-linear models, and when the percentage of missing data is high. This imputation technique distorts the original variable distribution.

1.4.4 How to do it...

```
1  # Replacing missing values with an arbitrary number
2
3  import pandas as pd
4  from sklearn.model_selection import train_test_split
5  from sklearn.impute import SimpleImputer
6  from feature_engine.imputation import ArbitraryNumberImputer
7
8  data = pd.read_csv('../data/credit_approval_uci.csv')
9
10 X_train, X_test, y_train, y_test = train_test_split(data.drop('target', axis='columns'),
11                                                    data['target'],
12                                                    test_size=.3,
13                                                    random_state=37)
14
15 arbitrary_cols = ['A2', 'A3', 'A8', 'A11']
16 print(X_train[arbitrary_cols].max().max())
17
18 X_train[arbitrary_cols] = X_train[arbitrary_cols].fillna(99)
19 X_test[arbitrary_cols] = X_test[arbitrary_cols].fillna(99)
20 print(X_train[arbitrary_cols].isnull().any().any())
21 print(X_test[arbitrary_cols].isnull().any().any())
22
```

```

23  # using scikit-learn
24  X_train, X_test, y_train, y_test = train_test_split(data.drop('target', axis='columns'),
25                                                    data['target'],
26                                                    test_size=.3,
27                                                    random_state=37)
28  imputer = SimpleImputer(strategy='constant', fill_value=99)
29  imputer.fit(X_train[arbitrary_cols])
30  X_train[arbitrary_cols] = imputer.transform(X_train[arbitrary_cols])
31  X_test[arbitrary_cols] = imputer.transform(X_test[arbitrary_cols])
32  print(X_train[arbitrary_cols].isnull().any().any())
33  print(X_test[arbitrary_cols].isnull().any().any())
34
35
36  # using feature-engine
37  X_train, X_test, y_train, y_test = train_test_split(data.drop('target', axis='columns'),
38                                                    data['target'],
39                                                    test_size=.3,
40                                                    random_state=37)
41  imputer = ArbitraryNumberImputer(arbitrary_number=99, variables=arbitrary_cols)
42
43  X_train = imputer.fit_transform(X_train)
44  X_test = imputer.transform(X_test)
45  print(X_train[arbitrary_cols].isnull().any().any())
46  print(X_test[arbitrary_cols].isnull().any().any())

```

第16行，首先要找到需要处理缺失值的最大值，然后将填充的任意值设置为比这个最大值要大。

第18-19行，也可以使用`imputation_dict`为每列单独设置填充值，将字典的key和value分别设置为字段名和想要填充的任意值。

第28行，如果你的数据集存在有缺失值的分类变量，那么`SimpleImputer()`也会用value将其填充。

On line 41, `ArbitraryNumberImputer()` can automatically select all numerical variables in the train set if we set the `variables` parameter to `None`.

1.5 Finding extreme values for imputation

Replacing missing values with a value at the end of the variable distribution (extreme values) is equivalent to replacing them with an arbitrary value, but instead of identifying the arbitrary values manually, these values are automatically selected as those at the very end of the variable distribution. Missing data can be replaced with a value that is greater or smaller than the remaining values in the variable. To select a value that is greater, we can use the mean plus a factor of the standard deviation, or the 75th quantile + (IQR * 1.5), where IQR is the

IQR given by the 75th quantile - the 25th quantile. To replace missing data with values that are smaller than the remaining values, we can use the mean minus a factor of the standard deviation, or the 25th quantile – (IQR * 1.5).

Tip

End-of-tail imputation may distort the distribution of the original variables, so it may not be suitable for linear models.

1.5.1 How to do it...

```
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from feature_engine.imputation import EndTailImputer
4
5 data = pd.read_csv('../data/credit_approval_uci.csv')
6
7 numeric_vars = [
8     var for var in data.select_dtypes(exclude='O').columns.to_list() if var != 'target'
9 ]
10 print(numeric_vars)
11
12 X_train, X_test, y_train, y_test = train_test_split(
13     data[numeric_vars],
14     data['target'],
15     test_size=.3,
16     random_state=0)
17
18 IQR = X_train.quantile(0.75) - X_train.quantile(0.25)
19 print(IQR)
20
21 imputation_dict = (X_train.quantile(.75) + 1.5 * IQR).to_dict()
22 print(imputation_dict)
23
24 X_train = X_train.fillna(value=imputation_dict)
25 X_test = X_test.fillna(value=imputation_dict)
26
27 print(X_train.isnull().any().any())
28 print(X_test.isnull().any().any())
```

```

29
30 # We can also replace missing data with values at the left tail of the distribution
31 X_train, X_test, y_train, y_test = train_test_split(
32     data[numeric_vars],
33     data['target'],
34     test_size=.3,
35     random_state=0)
36
37 imputer = EndTailImputer(
38     imputation_method='iqr',
39     tail='right',
40     fold=3,
41     variables=None)
42 imputer.fit(X_train)
43
44 print(imputer.imputer_dict_)
45
46 X_train = imputer.transform(X_train)
47 X_test = imputer.transform(X_test)
48
49 print(X_train.isnull().any().any())
50 print(X_test.isnull().any().any())

```

1.5.2 How it works...

On line 21, If we want to use the Gaussian approximation instead of the IQR proximity rule, we can calculate the value to replace missing data using `imputation_dict = (X_train.mean() + 3 * X_train.std()).to_dict()`.

On line 37-41, To use the mean and standard deviation to calculate the replacement values, we need to set `imputation_method="Gaussian"`. We can use 'left' or 'right' in the tail argument to specify the side of the distribution where we'll place the missing values.

1.6 Marking imputed values

A missing indicator is a binary variable that takes the value 1 or True to indicate whether a value was missing, or 0 or False otherwise. It is common practice to replace missing observations with the mean, median, or most frequent category while simultaneously marking those missing observations with missing indicators.

在数据库中，使用null查询会降低查询的性能，因此常见的操作是将缺失值转化为0来进行处理。

1.6.1 How to do it...

```
1  import pandas as pd
2  import numpy as np
3  from sklearn.model_selection import train_test_split
4  from sklearn.impute import SimpleImputer
5  from sklearn.compose import ColumnTransformer
6  from sklearn.pipeline import Pipeline
7  from feature_engine.imputation import (
8      AddMissingIndicator,
9      CategoricalImputer,
10     MeanMedianImputer,
11 )
12
13 data = pd.read_csv('../data/credit_approval_uci.csv')
14
15 X_train, X_test, y_train, y_test = train_test_split(
16     data.drop('target', axis='columns'), data['target'],
17     test_size=.3, random_state=37
18 )
19
20 varnames = ['A1', 'A3', 'A4', 'A5', 'A6', 'A7', 'A8']
21
22 indicators = [f'{var}_na' for var in varnames]
23
24 X_train[indicators] = X_train[varnames].isnull().astype(int)
25 X_test[indicators] = X_test[varnames].isnull().astype(int)
26
27 X_train.sample(5)
28
29 # using feature-engine
30 X_train, X_test, y_train, y_test = train_test_split(
31     data.drop('target', axis='columns'), data['target'],
32     test_size=.3, random_state=37
33 )
34
35 imputer = AddMissingIndicator(
36     variables=None,
37     missing_only=True,
```

```
38 )
39 imputer.fit(X_train)
40 imputer.variables_
41
42 X_train = imputer.transform(X_train)
43 X_test = imputer.transform(X_test)
44
45 pipe = Pipeline([
46     ('ind', AddMissingIndicator(missing_only=True)),
47     ('cat', CategoricalImputer(imputation_method='frequent')),
48     ('num', MeanMedianImputer(imputation_method='mean'))
49 ])
50
51 X_train = pipe.fit_transform(X_train)
52 X_test = pipe.fit_transform(X_test)
53
54 # using scikit-learn
55 X_train, X_test, y_train, y_test = train_test_split(
56     data.drop('target', axis='columns'), data['target'],
57     test_size=.3, random_state=37
58 )
59
60 num_vars = X_train.select_dtypes(exclude='O').columns.to_list()
61 cat_vars = X_train.select_dtypes(include='O').columns.to_list()
62
63 pipe = ColumnTransformer(
64     [
65         ('num_imputer',
66          SimpleImputer(strategy='mean',
67                        add_indicator=True),
68          num_vars),
69         ('cat_imputer',
70          SimpleImputer(strategy='most_frequent',
71                        add_indicator=True),
72          cat_vars)]
73 )
74 X_train = pipe.fit_transform(X_train)
75 X_test = pipe.fit_transform(X_test)
```


这段代码中，仔细看Pipeline的操作。

On line 45, 此处的Pipeline首先在DataFrame后面添加了存在missing value的字段，字段名设置为原字段名_na，然后将分类型变量设置用众数填充，数值型变量用均值填充。Feature-engine imputers automatically identify all numerical or categorical variables, modifying only the appropriate variables. So there is no need to slice the data or pass the variable names as arguments to the transformers in this case

1.7 Performing multivariate imputation by chained equations